

Predict Customer Conversion (Churn) with Machine Learning

Importing necessary libraries

```
import numpy as np
import pandas as pd
import sklearn
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

Reading and exploring the dataset

```
df = pd.read_csv("customer-churn-dataset.csv")
df
```

| customerID | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService |
|------------|--------|---------------|---------|------------|--------|--------------|
|------------|--------|---------------|---------|------------|--------|--------------|

```
df.shape
```

```
(7043, 21)
```

```
df.columns
```

```
df.columns.values
```

```
array(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
      'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract',
      'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges',
      'TotalCharges', 'Churn'], dtype=object)
```

```
...
```

```
...
```

```
...
```

```
...
```

```
...
```

```
...
```

```
...
```

```
...
```

To check for missing values or (NA)

```
df.isna().sum()
```

```
df.isna().sum()
```

```
customerID      0
gender          0
SeniorCitizen   0
Partner         0
Dependents      0
tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64
```

Dataset Statistics

```
df.describe()
```

| | SeniorCitizen | tenure | MonthlyCharges |
|--------------|---------------|-------------|----------------|
| count | 7043.000000 | 7043.000000 | 7043.000000 |
| mean | 0.162147 | 32.371149 | 64.761692 |
| std | 0.368612 | 24.559481 | 30.090047 |
| min | 0.000000 | 0.000000 | 18.250000 |
| max | 1.000000 | 56.000000 | 99.500000 |


```
df['Churn'].value_counts()
```

```
No      5174
Yes     1869
Name: Churn, dtype: int64
```

Visualizing the conversion

```
sns.countplot(df['Churn'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pas
FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7fdae08d7a50>
```



Percentage-wise results

```
numRetained = df[df.Churn == 'No'].shape[0]
numChurned = df[df.Churn == 'Yes'].shape[0]


# print the percentage of customers that stayed
print(numRetained/(numRetained + numChurned) * 100, '% of customers stayed with the company')
# print the percentage of customers that left
print(numChurned/(numRetained + numChurned) * 100, '% of customers left the company')
```

```
73.4630129206304 % of customers stayed with the company
26.536987079369588 % of customers left the company
```

Gender-wise visualization of customer conversion

```
sns.countplot(x='gender', hue='Churn', data=df)
```


<matplotlib.axes._subplots.AxesSubplot at 0x7fdadf621ad0>



Visualization of customer conversion for the internet service

```
sns.countplot(x='InternetService', hue='Churn', data=df)
```



<matplotlib.axes._subplots.AxesSubplot at 0x7fdadf15ba50>



Visualization of Numerical data

```
numericFeatures = ['tenure', 'MonthlyCharges']
fig, ax = plt.subplots(1,2, figsize=(28, 8))
df[df.Churn == "No"][numericFeatures].hist(bins=20, color='blue', alpha=0.5, ax=ax)
df[df.Churn == "Yes"][numericFeatures].hist(bins=20, color='orange', alpha=0.5, ax=ax)
```

```
array([<matplotlib.axes._subplots.AxesSubplot object at 0x7fdadf19c490>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7fdadf098a50>],
      dtype=object)
```



Dropping unnecessary columns from the dataset

```
cleanDF = df.drop('customerID', axis=1)
```

```
# Convert all the non-numeric columns to numeric
for column in cleanDF.columns:
    if cleanDF[column].dtype == np.number:
        continue
    cleanDF[column] = LabelEncoder().fit_transform(cleanDF[column])
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3: DeprecationWarning:
This is separate from the ipykernel package so we can avoid doing imports until

```
cleanDF.dtypes
```

```
gender          int64
SeniorCitizen   int64
Partner         int64
Dependents      int64
tenure          int64
PhoneService    int64
MultipleLines   int64
InternetService int64
OnlineSecurity  int64
OnlineBackup    int64
DeviceProtection int64
TechSupport     int64
StreamingTV     int64
StreamingMovies int64
Contract        int64
PaperlessBilling int64
PaymentMethod   int64
MonthlyCharges  float64
TotalCharges    int64
Churn           int64
dtype: object
```

Scaling of data

```
x = cleanDF.drop('Churn', axis=1)
y = cleanDF['Churn']
x = StandardScaler().fit_transform(x)
```

Split the data into 80% for training and 20% for testing

```
xtrain, xtest, ytrain, ytest = train_test_split(x,y, test_size=0.2, random_state=42)
```

Creating and Training the Logistic Regression Model

```
model = LogisticRegression()
# Train the model
model.fit(xtrain, ytrain)
```

```
LogisticRegression()
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                    warm_start=False)
```

```
LogisticRegression()
```

Creating predictions on the Test data

```
predictions = model.predict(xtest)

# print the predictions
print(predictions)
```

```
[1 0 0 ... 0 0 0]
```

Final scores - precision, recall and f1-score

```
print(classification_report(ytest, predictions))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.85 | 0.91 | 0.88 | 1036 |
| 1 | 0.69 | 0.56 | 0.62 | 373 |
| accuracy | | | 0.82 | 1409 |
| macro avg | 0.77 | 0.74 | 0.75 | 1409 |
| weighted avg | 0.81 | 0.82 | 0.81 | 1409 |